

Concepts

Events

An event is a set of values associated with a timestamp. It is a single entry of data and can have one or multiple lines. An event can be a text document, a configuration file, an entire stack trace, and so on. This is an example of an event in a web activity log:

```
173.26.34.223 - - [01/
Mar/2015:12:05:27 -0700] "GET /
trade/app?action=logout HTTP/1.1"
200 2953
```

You can also define transactions to search for and group together events that are conceptually related but span a duration of time. Transactions can represent a multistep business-related activity, such as all events related to a single customer session on a retail website.

Host, Source, and Source Type

A *host* is the name of the physical or virtual device where an event originates. The *host* field provides an easy way to find all data originating from a specific device. A *source* is the name of the file, directory, data stream, or other input from which a particular event originates. Sources are classified into *source types*, which can be either well known formats or formats defined by the user. Some common source types are HTTP web server logs and Windows event logs.

Events with the same source types can come from different sources. For example, events from the file `source=/var/log/messages` and from a syslog input port `source=UDP:514` often share the source type, `sourcetype=linux_syslog`.

Fields

Fields are searchable name and value pairings that distinguish one event from another. Not all events have the same fields and field values. Using fields, you can write tailored searches to retrieve the specific events that you want. When Splunk software processes events at index-time and search-time, the software extracts fields based on configuration file definitions and user-defined patterns.

Use the Field Extractor tool to automatically generate and validate field extractions at search-time. Regular expressions are automatically generated to extract fields. You can extract fields from events where values are separated by spaces, commas, or other characters.

Tags

A *tag* is a knowledge object that enables you to search for events that contain particular field values. You can assign one or more tags to any field/value combination, including event types, hosts, sources, and source types. Use tags to group related field values together, or to track abstract field values such as IP addresses or

ID numbers by giving them more descriptive names.

Events that match a specified search string can be saved as event types. Tag event types to organize your data into categories.

Index-Time and Search Time

During *index-time* processing, data is read from a source on a host and is classified into a source type. Timestamps are extracted, and the data is parsed into individual events. Line-breaking rules are applied to segment the events to display in the search results. Each event is written to an index on disk, where the event is later retrieved with a search request.

When a *search* starts, referred to as *search-time*, indexed events are retrieved from disk. *Fields* are extracted from the raw text for the event.

Indexes

When data is added, Splunk software parses the data into individual events, extracts the timestamp, applies line-breaking rules, and stores the events in an *index*. You can create new indexes for different inputs. By default, data is stored in the "main" index. Events are retrieved from one or more indexes during a search.

Core Features

Search

Search is the primary way users navigate data in Splunk software. You can write a search to retrieve events from an index, use statistical commands to calculate metrics and generate reports, search for specific conditions within a rolling time window, identify patterns in your data, predict future trends, and so on. You transform the events using the Splunk Search Process Language (SPL™). Searches can be saved as reports and used to power dashboards.

Reports

Reports are saved searches and pivots. You can run reports on an ad hoc basis, schedule reports to run on a regular interval, or set a scheduled report to generate alerts when the results meet particular conditions. Reports can be added to dashboards as dashboard panels.

Dashboards

Dashboards are made up of panels that contain modules such as search boxes, fields, and data visualizations. Dashboard panels are usually connected to saved searches or pivots. They can display the results of completed searches, as well as data from real-time searches.

Alerts

Alerts are triggered when search results meet specific conditions. You can use alerts on historical and real-time searches. Alerts can be configured to trigger actions such as sending alert information to designated email addresses or posting alert information to a web resource.

Additional Features (Splunk Enterprise only)

Data Model

A *data model* is a hierarchically-organized collection of datasets that Pivot uses to generate reports. Data model objects represent individual datasets, which the data model is composed of.

Pivot

Pivot refers to the table, chart, or other visualization you create using the Pivot Editor. You can map attributes defined by data model objects to data visualizations, without manually writing the searches. Pivots can be saved as reports and used to power dashboards.

Apps

Apps are a collection of configurations, knowledge objects, and customer designed views and dashboards. Apps extend the Splunk environment to fit the specific needs of organizational teams such as Unix or Windows system administrators, network security specialists, website managers, business analysts, and so on. A single Splunk Enterprise or Splunk Cloud installation can run multiple apps simultaneously.

Distributed Search

A *distributed search* provides a way to scale your deployment by separating the search management and presentation layer from the indexing and search retrieval layer. You use distributed search to facilitate horizontal scaling for enhanced performance, to control access to indexed data, and to manage geographically dispersed data.

Splunk Components

Forwarders

A Splunk instance that forwards data to another Splunk instance is referred to as a forwarder.

Indexer

An indexer is the Splunk instance that indexes data. The indexer transforms the raw data into events and stores the events into an index. The indexer also searches the indexed data in response to search requests. The search peers are indexers that fulfill search requests from the search head.

Search Head

In a distributed search environment, the search head is the Splunk instance that directs search requests to a set of search peers and merges the results back to the user. If the instance does only search and not indexing, it is usually referred to as a dedicated search head.

Search Processing Language

A Splunk search is a series of commands and arguments. Commands are chained together with a pipe "|" character to indicate that the output of one command feeds into the next command on the right.

```
search | command1 arguments1 |
command2 arguments2 | ...
```

At the start of the search pipeline, is an implied search command to retrieve events from the index. Search requests are written with keywords, quoted phrases, Boolean expressions, wildcards, field name/value pairs, and comparison expressions. The AND operator is implied between search terms. For example:

```
sourcetype=access_combined error |
top 5 uri
```

This search retrieves indexed web activity events that contain the term "error". For those events, it returns the top 5 most common URI values.

Search commands are used to filter unwanted events, extract more information, calculate values, transform, and statistically analyze the indexed data. Think of the search results retrieved from the index as a dynamically created table. Each indexed event is a row. The field values are columns. Each search command redefines the shape of that table. For example, search commands that filter events will remove rows, search commands that extract fields will add columns.

Time Modifiers

You can specify a time range to retrieve events inline with your search by using the `latest` and `earliest` search modifiers. The relative times are specified with a string of characters to indicate the amount of time (integer and unit) and an optional "snap to" time unit. The syntax is:

```
[+|-]<integer><unit>[@<snap _ time _
unit>]
```

The search `"error earliest=-1d@d latest=-h@h"` retrieves events containing "error" that occurred yesterday snapping to the beginning of the day (00:00:00) and through to the most recent hour of today, snapping on the hour.

The snap to time unit rounds the time down. For example, if it is 11:59:00 and you snap to hours (@h), the time used is 11:00:00 not 12:00:00. You can also snap to specific days of the week using @w0 for Sunday, @w1 for Monday, and so on.

Subsearches

A subsearch runs its own search and returns the results to the parent command as the argument value. The subsearch is run first and is contained in square brackets. For example, the following search uses a subsearch to find all syslog events from the user that had the last login error:

```
sourcetype=syslog [ search login
error | return 1 user ]
```

Optimizing Searches

The key to fast searching is to limit the data that needs to be pulled off disk to an absolute minimum. Then filter that data as early as possible in the search so that processing is done on the minimum data necessary.

Partition data into separate indexes, if you will rarely perform searches across multiple types of data. For example, put web data in one index, and firewall data in another.

Limit the time range to only what is needed. For example `-1h not -1w`, or `earliest=-1d`.

Search as specifically as you can. For example, `fatal_error not *error*`

Filter out results as soon as possible before calculations. Use field-value pairs, before the first pipe. For example, `>ERROR status=404 |...` instead of `>ERROR | search status=404...` Or use filtering commands such as `where`.

Filter out unnecessary fields as soon as possible in the search.

Postpone commands that process over the entire result set (non-streaming commands) as late as possible in your search. Some of these commands are: `dedup`, `sort`, and `stats`.

Use post-processing searches in dashboards.

Use summary indexing, and report and data model acceleration features.

Machine Learning Commands

The Machine Learning Toolkit delivers additional SPL commands that you can use to apply machine learning to your data. Find out more in the Machine Learning Quick Reference Guide.

Common Search Commands

Command	Description
chart/ timechart	Returns results in a tabular output for (time-series) charting.
dedup	Removes subsequent results that match a specified criterion.
eval	Calculates an expression. See COMMON EVAL FUNCTIONS.
fields	Removes fields from search results.
head/tail	Returns the first/last N results.
lookup	Adds field values from an external source.
rename	Renames a field. Use wildcards to specify multiple fields.
rex	Specifies regular expression named groups to extract fields.
search	Filters results to those that match the search expression.
sort	Sorts the search results by the specified fields.
stats	Provides statistics, grouped optionally by fields. See COMMON STATS FUNCTIONS.
table	Specifies fields to keep in the result set. Retains data in tabular format.
top/rare	Displays the most/least common values of a field.
transaction	Groups search results into transactions.
where	Filters search results using eval expressions. Used to compare two different fields.

splunk>

www.splunk.com
docs.splunk.com

Splunk Inc.
250 Brannan Street
San Francisco, CA 94107

Copyright © 2016 Splunk Inc. All rights reserved. Splunk, Splunk>, Listen to Your Data, The Engine for Machine Data, Hunk, Splunk Cloud, Splunk Light, SPL and Splunk MINT are trademarks and registered trademarks of Splunk Inc. in the United States and other countries. All other brand names, product names, or trademarks belong to their respective owners.

Common Eval Functions

The eval command calculates an expression and puts the resulting value into a field (e.g. "...| eval force = mass * acceleration"). The following table lists some of the functions used with the eval command. You can also use basic arithmetic operators (+ - * / %), string concatenation (e.g., "...| eval name = last . "," . first"), and Boolean operations (AND OR NOT XOR < > <= >= != == LIKE).

Function	Description	Examples
abs(X)	Returns the absolute value of X.	abs(number)
case(X,"Y",...)	Takes pairs of arguments X and Y, where X arguments are Boolean expressions. When evaluated to TRUE, the arguments return the corresponding Y argument.	case(error == 404, "Not found", error == 500, "Internal Server Error", error == 200, "OK")
ceil(X)	Ceiling of a number X.	ceil(1.9)
cidrmatch("X",Y)	Identifies IP addresses that belong to a particular subnet.	cidrmatch("123.132.32.0/25",ip)
coalesce(X,...)	Returns the first value that is not null.	coalesce(null(), "Returned val", null())
cos(X)	Calculates the cosine of X.	n=cos(0)
exact(X)	Evaluates an expression X using double precision floating point arithmetic.	exact(3.14*num)
exp(X)	Returns eX.	exp(3)
if(X,Y,Z)	If X evaluates to TRUE, the result is the second argument Y. If X evaluates to FALSE, the result evaluates to the third argument Z.	if(error==200, "OK", "Error")
isbool(X)	Returns TRUE if X is Boolean.	isbool(field)
isint(X)	Returns TRUE if X is an integer.	isint(field)
isnull(X)	Returns TRUE if X is NULL.	isnull(field)
isstr()	Returns TRUE if X is a string.	isstr(field)
len(X)	This function returns the character length of a string X.	len(field)
like(X,"Y")	Returns TRUE if and only if X is like the SQLite pattern in Y.	like(field, "addr%")
log(X,Y)	Returns the log of the first argument X using the second argument Y as the base. Y defaults to 10.	log(number,2)
lower(X)	Returns the lowercase of X.	lower(username)
ltrim(X,Y)	Returns X with the characters in Y trimmed from the left side. Y defaults to spaces and tabs.	ltrim(" ZZZabcZZ ", " Z")
match(X,Y)	Returns if X matches the regex pattern Y.	match(field, "^\\d{1,3}\\\\.\\d\$")
max(X,...)	Returns the maximum.	max(delay, mydelay)
md5(X)	Returns the MD5 hash of a string value X.	md5(field)
min(X,...)	Returns the minimum.	min(delay, mydelay)
mvcount(X)	Returns the number of values of X.	mvcount(multifield)
mvfilter(X)	Filters a multi-valued field based on the Boolean expression X.	mvfilter(match(email, "net\$"))
mvindex(X,Y,Z)	Returns a subset of the multivalued field X from start position (zero-based) Y to Z (optional).	mvindex(multifield, 2)
mvjoin(X,Y)	Given a multi-valued field X and string delimiter Y, and joins the individual values of X using Y.	mvjoin(address, ",")
now()	Returns the current time, represented in Unix time.	now()
null()	This function takes no arguments and returns NULL.	null()
nullif(X,Y)	Given two arguments, fields X and Y, and returns the X if the arguments are different. Otherwise returns NULL.	nullif(fieldA, fieldB)
random()	Returns a pseudo-random number ranging from 0 to 2147483647.	random()
relative_time(X,Y)	Given epochtime time X and relative time specifier Y, returns the epochtime value of Y applied to X.	relative_time(now(),"-1d@d")
replace(X,Y,Z)	Returns a string formed by substituting string Z for every occurrence of regex string Y in string X.	Returns date with the month and day numbers switched, so if the input was 4/30/2015 the return value would be 30/4/2009: replace(date, "\\d{1,2})/\\d{1,2})/", "\\2/\\1/")

Common Eval Functions (continued)

Function	Description	Examples
round(X,Y)	Returns X rounded to the amount of decimal places specified by Y. The default is to round to an integer.	<code>round(3.5)</code>
rtrim(X,Y)	Returns X with the characters in Y trimmed from the right side. If Y is not specified, spaces and tabs are trimmed.	<code>rtrim(" ZZZZabcZZ ", " Z")</code>
searchmatch(X)	Returns true if the event matches the search string X.	<code>searchmatch("foo AND bar")</code>
split(X,"Y")	Returns X as a multi-valued field, split by delimiter Y.	<code>split(address, ";")</code>
sqr(X)	Returns the square root of X.	<code>sqr(9)</code>
strftime(X,Y)	Returns epochtime value X rendered using the format specified by Y.	<code>strftime(_time, "%H:%M")</code>
strptime(X,Y)	Given a time represented by a string X, returns value parsed from format Y.	<code>strptime(timeStr, "%H:%M")</code>
substr(X,Y,Z)	Returns a substring field X from start position (1-based) Y for Z (optional) characters.	<code>substr("string", 1, 3)</code>
time()	Returns the wall-clock time with microsecond resolution.	<code>time()</code>
tonumber(X,Y)	Converts input string X to a number, where Y (optional, defaults to 10) defines the base of the number to convert to.	<code>tonumber("0A4",16)</code>
tostring(X,Y)	Returns a field value of X as a string. If the value of X is a number, it reformats it as a string. If X is a Boolean value, reformats to "True" or "False". If X is a number, the second argument Y is optional and can either be "hex" (convert X to hexadecimal), "commas" (formats X with commas and 2 decimal places), or "duration" (converts seconds X to readable time format HH:MM:SS).	This example returns: <code>foo=615</code> and <code>foo2=00:10:15</code> : ... eval foo=615 eval foo2 = tostring(foo, "duration")
typeof(X)	Returns a string representation of the field type.	This example returns: "NumberStringBoolInvalid": <code>typeof(12)+typeof("string")</code>
urldecode(X)	Returns the URL X decoded.	<code>urldecode("http%3A%2F%2Fwww.splunk.com%2Fdownload%3Fr%3Dheader")</code>
validate(X,Y,...)	Given pairs of arguments, Boolean expressions X and strings Y, returns the string Y corresponding to the first expression X that evaluates to False and defaults to NULL if all are True.	<code>validate(isint(port), "ERROR: Port is not an integer", port >= 1 AND port <= 65535, "ERROR: Port is out of range")</code>

Common Stats Functions

Common statistical functions used with the chart, stats, and timechart commands. Field names can be wildcarded, so `avg(*delay)` might calculate the average of the delay and xdelay fields.

avg(X)	Returns the average of the values of field X.
count(X)	Returns the number of occurrences of the field X. To indicate a specific field value to match, format X as <code>eval(field="value")</code> .
dc(X)	Returns the count of distinct values of the field X.
earliest(X)	Returns the chronologically earliest seen value of X.
latest(X)	Returns the chronologically latest seen value of X.
max(X)	Returns the maximum value of the field X. If the values of X are non-numeric, the max is found from alphabetical ordering.
median(X)	Returns the middle-most value of the field X.
min(X)	Returns the minimum value of the field X. If the values of X are non-numeric, the min is found from alphabetical ordering.
mode(X)	Returns the most frequent value of the field X.
perc<X>(Y)	Returns the X-th percentile value of the field Y. For example, <code>perc5(total)</code> returns the 5th percentile value of a field "total".
range(X)	Returns the difference between the max and min values of the field X.
stdev(X)	Returns the sample standard deviation of the field X.
stdevp(X)	Returns the population standard deviation of the field X.
sum(X)	Returns the sum of the values of the field X.
sumsq(X)	Returns the sum of the squares of the values of the field X.
values(X)	Returns the list of all distinct values of the field X as a multi-value entry. The order of the values is alphabetical.
var(X)	Returns the sample variance of the field X.

Search Examples

Filter Results

Returns X rounded to the amount of decimal places specified by Y. The default is to round to an integer.	<code>round(3.5)</code>
Returns X with the characters in Y trimmed from the right side. If Y is not specified, spaces and tabs are trimmed.	<code>rtrim(" ZZZZabcZZ ", "Z")</code>
Returns true if the event matches the search string X.	<code>searchmatch("foo AND bar")</code>
Returns X as a multi-valued field, split by delimiter Y.	<code>split(address, ";")</code>
Given pairs of arguments, Boolean expressions X and strings Y, returns the string Y corresponding to the first expression X that evaluates to False and defaults to NULL if all are True.	<code>validate(isint(port), "ERROR: Port is not an integer", port >= 1 AND port <= 65535, "ERROR: Port is out of range")</code>

Group Results

Cluster results together, sort by their "cluster_count" values, and then return the 20 largest clusters (in data size).	<code>... cluster t=0.9 showcount=true sort limit=20 -cluster_count</code>
Group results that have the same "host" and "cookie", occur within 30 seconds of each other, and do not have a pause greater than 5 seconds between each event into a transaction.	<code>... transaction host cookie maxspan=30s maxpause=5s</code>
Group results with the same IP address (clientip) and where the first result contains "signon", and the last result contains "purchase".	<code>... transaction clientip startswith="signon" endswith="purchase"</code>

Order Results

Return the first 20 results.	<code>... head 20</code>
Reverse the order of a result set.	<code>... reverse</code>
Sort results by "ip" value (in ascending order) and then by "url" value (in descending order).	<code>... sort ip, -url</code>
Return the last 20 results in reverse order.	<code>... tail 20</code>

Reporting

Return the maximum "delay" by "size", where "size" is broken down into a maximum of 10 equal sized buckets.	<code>... chart max(delay) by size bins=10</code>
Return max(delay) for each value of foo split by the value of bar.	<code>... chart max(delay) over foo by bar</code>
Return max(delay) for each value of foo.	<code>... chart max(delay) over foo</code>
Count the events by "host"	<code>... stats count by host</code>

Reporting (cont.)

Create a table showing the count of events and a small line chart	<code>... stats sparkline count by host</code>
Create a timechart of the count of from "web" sources by "host"	<code>... timechart count by host</code>
Calculate the average value of "CPU" each minute for each "host".	<code>... timechart span=1m avg(CPU) by host</code>
Return the average for each hour, of any unique field that ends with the string "lay" (e.g., delay, xdelay, relay, etc).	<code>... stats avg(*lay) by date_hour</code>
Return the 20 most common values of the "url" field.	<code>... top limit=20 url</code>
Return the least common values of the "url" field.	<code>... rare url</code>

Advanced Reporting

Compute the overall average duration and add 'avgdur' as a new field to each event where the 'duration' field exists	<code>... eventstats avg(duration) as avgdur</code>
Find the cumulative sum of bytes.	<code>... streamstats sum(bytes) as bytes_total timechart max(bytes_total)</code>
Find anomalies in the field 'Close_Price' during the last 10 years.	<code>sourcetype=nasdaq earliest=-10y anomalydetection Close_Price</code>
Create a chart showing the count of events with a predicted value and range added to each event in the time-series.	<code>... timechart count predict count</code>
Computes a five event simple moving average for field 'count' and write to new field 'smoothed_count'.	<code>"... timechart count trendline sma5(count) as smoothed_count"</code>

Add Fields

Set velocity to distance / time.	<code>... eval velocity=distance/time</code>
Extract "from" and "to" fields using regular expressions. If a raw event contains "From: Susan To: David", then from=Susan and to=David.	<code>... rex field=_raw "From: (?<from>.*) To: (?<to>.*)"</code>
Save the running total of "count" in a field called "total_count".	<code>... accum count as total_count</code>
For each event where 'count' exists, compute the difference between count and its previous value and store the result in 'countdiff'.	<code>... delta count as countdiff</code>

Filter Fields

Keep only the "host" and "ip" fields, and display them in that order.	<code>... fields + host, ip</code>
Remove the "host" and "ip" fields from the results.	<code>... fields - host, ip</code>

Search Examples (continued)

Lookup Tables (Splunk Enterprise only)

For each event, use the lookup table usertogroup to locate the matching "user" value from the event. Output the group field value to the event	... lookup usertogroup user output group
Read in the usertogroup lookup table that is defined in the transforms.conf file.	... inputlookup usertogroup
Write the search results to the lookup file "users.csv".	... outputlookup users.csv

Modify Fields

Rename the "_ip" field as "IPAddress".	... rename _ip as IPAddress
--	-------------------------------

Regular Expressions (Regexes)

Regular Expressions are useful in multiple areas: search commands regex and rex; eval functions match() and replace(); and in field extraction.

Regex	Note	Example	Explanation
\s	white space	\d\s\d	digit space digit
\S	not white space	\d\S\d	digit non-whitespace digit
\d	digit	\d\d\d-\d\d-\d\d\d\d	SSN
\D	not digit	\D\D\D	three non-digits
\w	word character (letter, number, or _)	\w\w\w	three word chars
\W	not a word character	\W\W\W	three non-word chars
[...]	any included character	[a-z0-9#]	any char that is a thru z, 0 thru 9, or #
[^...]	no included character	[^xyz]	any char but x, y, or z
*	zero or more	\w*	zero or more words chars
+	one or more	\d+	integer
?	zero or one	\d\d\d-?\d\d-?\d\d\d\d	SSN with dashes being optional
	or	\w \d	word or digit character
(?P<var>...)	named extraction	(?P<ssn>\d\d\d-\d\d-\d\d\d\d)	pull out a SSN and assign to 'ssn' field
(?: ...)	logical or atomic grouping	(?:[a-zA-Z])\d	alphabetic character OR a digit
^	start of line	^d+	line begins with at least one digit
\$	end of line	\d+\$	line ends with at least one digit
{...}	number of repetitions	\d{3,5}	between 3-5 digits
\	escape	\[escape the [character

Multi-Valued Fields

Combine the multiple values of the recipients field into a single value	... nomv recipients
Separate the values of the "recipients" field into multiple field values, displaying the top recipients	... makemv delim="," recipients top recipients
Create new results for each value of the multivalued field "recipients"	... mvexpand recipients
Find the number of recipient values	... eval to_count = mvcount(recipients)
Find the first email address in the recipient field	... eval recipient_first = mvindex(recipient,0)
Find all recipient values that end in .net or .org	... eval netorg_recipients = mvfilter match(recipient, ".net\$") OR match(recipient, ".org\$")
Find the index of the first recipient value match ".org\$"	... eval orgindex = mvfind(recipient, ".org\$")

Common Date and Time Formatting

Use these values for eval functions strftime() and strptime(), and for timestamping event data.

Time	%H	24 hour (leading zeros) (00 to 23)
	%I	12 hour (leading zeros) (01 to 12)
	%M	Minute (00 to 59)
	%S	Second (00 to 61)
	%N	subseconds with width (%3N = millisecs, %6N = microsecs, %9N = nanosecs)
	%p	AM or PM
	%Z	Time zone (EST)
	%z	Time zone offset from UTC, in hour and minute: +hhmm or -hhmm. (-0500 for EST)
	%s	Seconds since 1/1/1970 (1308677092)
Days	%d	Day of month (leading zeros) (01 to 31)
	%j	Day of year (001 to 366)
	%w	Weekday (0 to 6)
	%a	Abbreviated weekday (Sun)
	%A	Weekday (Sunday)
Months	%b	Abbreviated month name (Jan)
	%B	Month name (January)
	%m	Month number (01 to 12)
Years	%y	Year without century (00 to 99)
	%Y	Year (2015)
Examples	%Y-%m-%d	2014-12-31
	%y-%m-%d	14-12-31
	%b %d, %Y	Jan 24, 2015
	%B %d, %Y	January 24, 2015
	q %d %b ' %y = %Y-%m-%d	q 25 Feb '15 = 2015-02-25

For more info visit:
docs.splunk.com